

ml-lab-02-aie23114

February 16, 2025

```
[2]: from google.colab import files
      uploaded = files.upload()
```

<IPython.core.display.HTML object>

Saving Lab Session Data.xlsx to Lab Session Data.xlsx

A1. Please refer to the “Purchase Data” worksheet of Lab Session Data.xlsx. Please load the data and segregate them into 2 matrices A & C (following the nomenclature of $AX = C$). Do the following activities. • What is the dimensionality of the vector space for this data? • How many vectors exist in this vector space? • What is the rank of Matrix A? • Using Pseudo-Inverse find the cost of each product available for sale.

(Suggestion: If you use Python, you can use `numpy.linalg.pinv()` function to get a pseudo-inverse.)

A2. Use the Pseudo-inverse to calculate the model vector X for predicting the cost of the products available with the vendor. A3. Mark all customers (in “Purchase Data” table) with payments above Rs. 200 as RICH and others as POOR. Develop a classifier model to categorize customers into RICH or POOR class based on purchase behavior.

```
[12]: import pandas as pd
      import numpy as np
      df = pd.read_excel(r'/content/Lab Session Data.xlsx', sheet_name="Purchase_
      ↪data")
      df = df.iloc[:, :5]
      print(df)

      A = df.iloc[:10, 1:4].values
      print(f"A:{A}")
      C = df.iloc[:10, 4].values.reshape(-1, 1)
      print(f"C:{C}")

      print(f"Dimensionality of the vector space:{df.shape}")

      print(f"Number of vectors:{df.shape[0]}")

      rank_A = np.linalg.matrix_rank(A)
      print(f"Rank of matrix A: {rank_A}")

      A_pinv = np.linalg.pinv(A)
      print(f"pinv of A is {A_pinv}")
```

```

X = np.dot(A_pinv, C)
print(f"model vector X is {X}")
print(f" Cost of each candy is Rs. {X[0]}")
print(f" Cost of each mango is Rs. {X[1]}")
print(f" Cost of each milk packet is Rs. {X[2]}")

category = []
for payment in df['Payment (Rs)']:
    if payment >= 200 :
        category.append('Rich')
    else:
        category.append('Poor')
df['category'] = category
print(df)

```

	Customer	Candies (#)	Mangoes (Kg)	Milk Packets (#)	Payment (Rs)
0	C_1	20	6	2	386
1	C_2	16	3	6	289
2	C_3	27	6	2	393
3	C_4	19	1	2	110
4	C_5	24	4	2	280
5	C_6	22	1	5	167
6	C_7	15	4	2	271
7	C_8	18	4	2	274
8	C_9	21	1	4	148
9	C_10	16	2	4	198

A: [[20 6 2]

[16 3 6]

[27 6 2]

[19 1 2]

[24 4 2]

[22 1 5]

[15 4 2]

[18 4 2]

[21 1 4]

[16 2 4]]

C: [[386]

[289]

[393]

[110]

[280]

[167]

[271]

[274]

[148]

[198]]

```

Dimensionality of the vector space:(10, 5)
Number of vectors:10
Rank of matrix A: 3
pinv of A is [[-0.01008596 -0.03124505  0.01013951  0.0290728  0.0182907
0.01161794
-0.00771348  0.00095458  0.01743623 -0.00542016]
[ 0.09059668  0.07263726  0.03172933 -0.09071908 -0.01893196 -0.06926996
 0.05675464  0.03152577 -0.07641966  0.00357352]
[ 0.00299878  0.15874243 -0.05795468 -0.06609024 -0.06295043  0.03348017
 0.01541831 -0.01070461  0.00029003  0.05938755]]
model vector X is [[ 1.]
[55.]
[18.]]
Cost of each candy is Rs. [1.]
Cost of each mango is Rs. [55.]
Cost of each milk packet is Rs. [18.]

```

	Customer	Candies (#)	Mangoes (Kg)	Milk Packets (#)	Payment (Rs)	category
0	C_1	20	6	2	386	Rich
1	C_2	16	3	6	289	Rich
2	C_3	27	6	2	393	Rich
3	C_4	19	1	2	110	Poor
4	C_5	24	4	2	280	Rich
5	C_6	22	1	5	167	Poor
6	C_7	15	4	2	271	Rich
7	C_8	18	4	2	274	Rich
8	C_9	21	1	4	148	Poor
9	C_10	16	2	4	198	Poor

A4. Please refer to the data present in “IRCTC Stock Price” data sheet of the above excel file. Do the following after loading the data to your programming platform. • Calculate the mean and variance of the Price data present in column D.

(Suggestion: if you use Python, you may use statistics.mean() & statistics.variance() methods).

- Select the price data for all Wednesdays and calculate the sample mean. Compare the mean with the population mean and note your observations.
- Select the price data for the month of Apr and calculate the sample mean. Compare the mean with the population mean and note your observations.
- From the Chg% (available in column I) find the probability of making a loss over the stock. (Suggestion: use lambda function to find negative values)
- Calculate the probability of making a profit on Wednesday.
- Calculate the conditional probability of making profit, given that today is Wednesday.
- Make a scatter plot of Chg% data against the day of the week

```

[13]: import pandas as pd
import numpy as np
import statistics
import matplotlib.pyplot as plt

df = pd.read_excel(r'/content/Lab Session Data.xlsx', sheet_name="IRCTC Stock_
↪Price")
df = df.iloc[:, :9]

```

```

print(df)

mean = statistics.mean(df["Price"])
var = statistics.variance(df["Price"])
print(f"Price Mean: {mean}")
print(f"Price Variance: {var}")

wed_mean = df.loc[df['Day'] == 'Wed', 'Price'].mean()
print(f"Mean Price on Wednesdays: {wed_mean}")
print("Observation: Sales at IRCTC are lower on Wednesdays compared to the_
↳overall average.")

apr_mean = df.loc[df['Month'] == 'Apr', 'Price'].mean()
print(f"Mean Price in April: {apr_mean}")
print("Observation: Sales at IRCTC are higher in April compared to the overall_
↳average.")

loss_probability = sum(df['Chg%'] < 0) / len(df)
print(f"Probability of making a loss: {loss_probability}")

wed_df = df[df['Day'] == 'Wed']
wed_profit = sum(wed_df['Chg%'] > 0) / len(wed_df)
print(f"Probability of making a profit on Wednesday: {wed_profit}")

profit_given_wed = (sum((df['Day'] == 'Wed') & (df['Chg%'] > 0)) /
↳sum(df['Day'] == 'Wed'))
print(f"Conditional Probability (Profit | Wednesday): {profit_given_wed}")

days = df['Day']
chg = df['Chg%']
plt.scatter(days, chg)
plt.xlabel("Day of the Week")
plt.ylabel("Chg%")
plt.title("Scatter Plot of chg% against day of the week")
plt.show()

```

	Date	Month	Day	Price	Open	High	Low	Volume	\
0	Jun 29, 2021	Jun	Tue	2081.85	2092.00	2126.90	2065.05	1.67M	
1	Jun 28, 2021	Jun	Mon	2077.75	2084.00	2112.45	2068.40	707.73K	
2	Jun 25, 2021	Jun	Fri	2068.85	2084.35	2088.50	2053.10	475.82K	
3	Jun 24, 2021	Jun	Thu	2072.95	2098.00	2098.00	2066.00	541.51K	
4	Jun 23, 2021	Jun	Wed	2078.25	2102.00	2111.40	2072.00	809.62K	
..	
244	Jul 07, 2020	Jul	Tue	1397.40	1410.00	1411.00	1390.05	480.21K	
245	Jul 06, 2020	Jul	Mon	1400.75	1405.50	1415.50	1394.00	614.93K	
246	Jul 03, 2020	Jul	Fri	1405.10	1415.00	1425.00	1398.00	599.49K	
247	Jul 02, 2020	Jul	Thu	1412.35	1440.00	1467.80	1395.30	2.16M	

248 Jul 01, 2020 Jul Wed 1363.05 1363.65 1377.00 1356.00 383.00K

Chg%
0 0.0020
1 0.0043
2 -0.0020
3 -0.0026
4 -0.0023
..
244 -0.0024
245 -0.0031
246 -0.0051
247 0.0362
248 0.0032

[249 rows x 9 columns]

Price Mean: 1560.663453815261

Price Variance: 58732.365352539186

Mean Price on Wednesdays: 1550.7060000000001

Observation: Sales at IRCTC are lower on Wednesdays compared to the overall average.

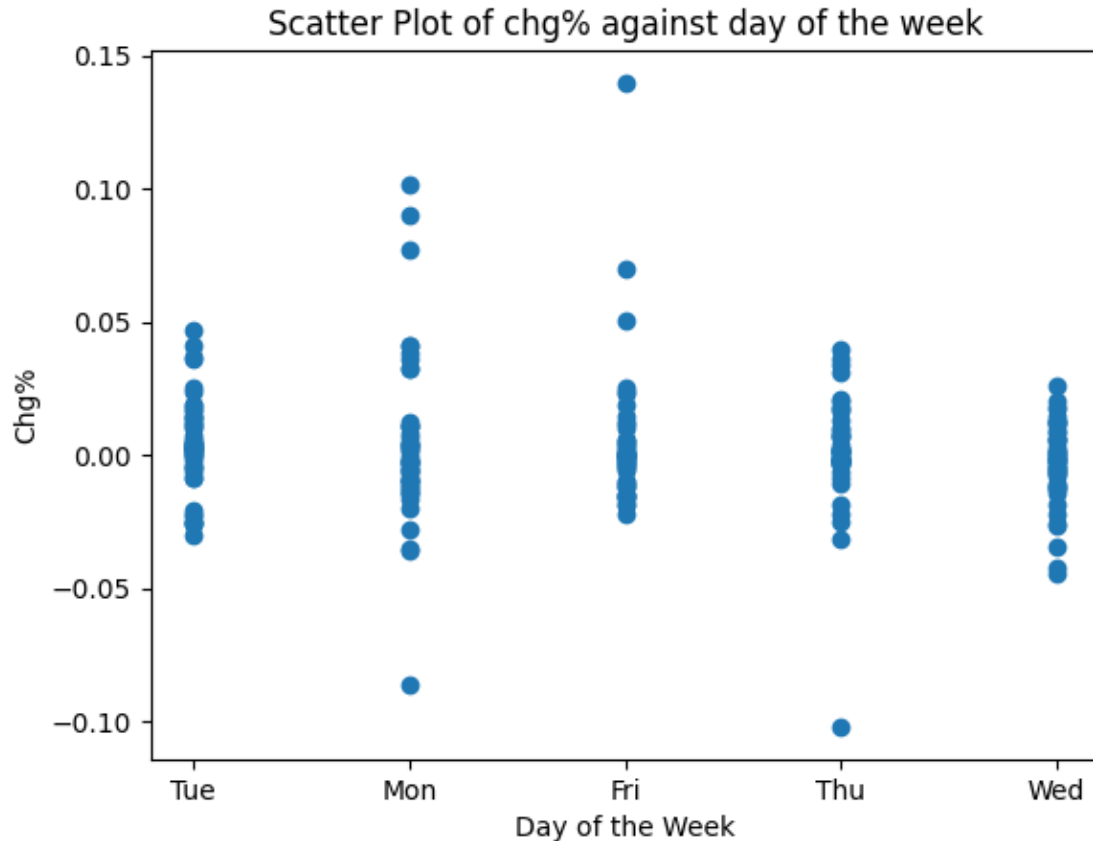
Mean Price in April: 1698.9526315789474

Observation: Sales at IRCTC are higher in April compared to the overall average.

Probability of making a loss: 0.4979919678714859

Probability of making a profit on Wednesday: 0.42

Conditional Probability (Profit | Wednesday): 0.42



A5. Data Exploration: Load the data available in “thyroid0387_UCI” worksheet. Perform the following tasks:

- Study each attribute and associated values present. Identify the datatype (nominal etc.) for the attribute.
- For categorical attributes, identify the encoding scheme to be employed. (Guidance: employ label encoding for ordinal variables while One-Hot encoding may be employed for nominal variables).
- Study the data range for numeric variables.
- Study the presence of missing values in each attribute.
- Study presence of outliers in data.

• For numeric variables, calculate the mean and variance (or standard deviation). A6. Data Imputation: employ appropriate central tendencies to fill the missing values in the data variables. Employ following guidance.

- Mean may be used when the attribute is numeric with no outliers

- Median may be employed for attributes which are numeric and contain outliers
- Mode may be employed for categorical attributes

A7. Data Normalization / Scaling: from the data study, identify the attributes which may need normalization. Employ appropriate normalization techniques to create normalized set of data.

A8. Similarity Measure: Take the first 2 observation vectors from the dataset. Consider only the attributes (direct or derived) with binary values for these vectors (ignore other attributes). Calculate the Jaccard Coefficient (JC) and Simple Matching Coefficient (SMC) between the document vectors. Use first vector for each document for this. Compare the values for JC and SMC and judge the appropriateness of each of them.

JC = $\frac{f_{11}}{f_{01} + f_{10} + f_{11}}$ SMC = $\frac{f_{11} + f_{00}}{f_{00} + f_{01} + f_{10} + f_{11}}$ f11= number of attributes where the attribute carries value of 1 in both the vectors.

A9. Cosine Similarity Measure: Now take the complete vectors for these two observations (including all the attributes). Calculate the Cosine similarity

between the documents by using the second feature vector for each document. A10. Heatmap Plot: Consider the first 20 observation vectors. Calculate the JC, SMC and COS between the pairs of vectors for these 20 vectors. Employ similar strategies for coefficient calculation as in A4 & A5. Employ a heatmap plot to visualize the similarities.

```
[20]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder, MinMaxScaler
from sklearn.metrics.pairwise import cosine_similarity

df = pd.read_excel(r"/content/Lab Session Data.xlsx",
    ↪sheet_name="thyroid0387_UCI")
# Replace '?' with NaN and infer proper datatypes
df.replace('?', np.nan, inplace=True)
df = df.infer_objects()

data_types = {}
# Identify and classify data types: Nominal, Ordinal, Ratio, Interval
for col in df.columns:
    if df[col].dtype == 'object':
        if df[col].nunique() <= 10:
            data_types[col] = 'Ordinal'
        else:
            data_types[col] = 'Nominal'
    elif df[col].dtype in ['int64', 'float64']:
        if df[col].min() >= 0:
            data_types[col] = 'Ratio'
        else:
            data_types[col] = 'Interval'
print("Attribute Data Types:")
for col, dtype in data_types.items():
    print(f"{col}: {dtype}")

# Identify categorical columns
categorical_cols = df.select_dtypes(include=['object']).columns

# Encoding Scheme
print("Encoding Scheme:")
for col in categorical_cols:
    if data_types[col] == 'Ordinal':
        print(f"{col}: Label Encoding (Ordinal)")
    else:
        print(f"{col}: One-Hot Encoding (Nominal)")

numerical_cols = df.select_dtypes(include=['float64', 'int64']).columns
```

```

print("Range for Numeric Variables:")
for col in numerical_cols:
    min_value = df[col].min()
    max_value = df[col].max()
    data_range = max_value - min_value
    print(f"{col}: Min = {min_value}, Max = {max_value}, Range = {data_range}")
print("Mean and Standard Deviation (or Variance) for Numeric Variables:")
for col in numerical_cols:
    mean_value = df[col].mean()
    std_dev = df[col].std()
    variance = std_dev ** 2
    print(f"{col}: Mean = {mean_value:.2f}, Standard Deviation = {std_dev},  

    ↪Variance = {variance}")

# Check for missing values
missing_values = df.isnull().sum()
print("Missing Values:", missing_values)

# Apply Label Encoding to categorical features
for col in categorical_cols:
    df[col] = df[col].astype(str)
    df[col] = LabelEncoder().fit_transform(df[col])

# Identify outliers
Q1 = df.quantile(0.25)
Q3 = df.quantile(0.75)
IQR = Q3 - Q1
outliers = ((df < (Q1 - 1.5 * IQR)) | (df > (Q3 + 1.5 * IQR))).sum()
print("Outliers per column:", outliers)

# Data Imputation
for col in df.columns:
    if df[col].dtype in ['float64', 'int64']:
        if outliers[col] > 0:
            df[col] = df[col].fillna(df[col].median())
        else:
            df[col] = df[col].fillna(df[col].mean())
    else:
        df[col] = df[col].fillna(df[col].mode()[0])
print("Missing values after imputation:", df.isnull().sum())

#Data Normalization
scaler = MinMaxScaler()
numerical_cols = df.select_dtypes(include=['float64', 'int64']).columns
df[numerical_cols] = scaler.fit_transform(df[numerical_cols])
print("Normalized Data Sample:", df.head())

```



```

# jc, smc
vector1 = df.iloc[0, :].values
vector2 = df.iloc[1, :].values
f11 = np.sum((vector1 == 1) & (vector2 == 1))
f00 = np.sum((vector1 == 0) & (vector2 == 0))
f10 = np.sum((vector1 == 1) & (vector2 == 0))
f01 = np.sum((vector1 == 0) & (vector2 == 1))
if (f01 + f10 + f11) != 0:
    jc = f11 / (f01 + f10 + f11)
else:
    jc = 0
if (f00 + f01 + f10 + f11) != 0:
    smc = (f11 + f00) / (f00 + f01 + f10 + f11)
else:
    smc = 0
print(f"Jaccard Coefficient: {jc}, SMC: {smc}")

# Cosine Similarity
vector1 = df.iloc[0, :].values.reshape(1, -1)
vector2 = df.iloc[1, :].values.reshape(1, -1)
cosine_sim = cosine_similarity(vector1, vector2)[0][0]
print(f"Cosine Similarity: {cosine_sim}")

# Heatmap plot
df_subset = df.iloc[:20, :]
similarity_matrix = np.zeros((20, 20))
for i in range(20):
    for j in range(20):
        if i != j:
            similarity_matrix[i, j] = np.linalg.norm(df_subset.iloc[i] -
↳df_subset.iloc[j])
sns.heatmap(similarity_matrix)
plt.show()

```

<ipython-input-20-cadb9acfc739>:10: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in a future version. To retain the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-in to the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`
df.replace('?', np.nan, inplace=True)

Attribute Data Types:

Record ID: Ratio

age: Ratio

sex: Ordinal

on thyroxine: Ordinal

query on thyroxine: Ordinal

on antithyroid medication: Ordinal

sick: Ordinal

pregnant: Ordinal
thyroid surgery: Ordinal
I131 treatment: Ordinal
query hypothyroid: Ordinal
query hyperthyroid: Ordinal
lithium: Ordinal
goitre: Ordinal
tumor: Ordinal
hypopituitary: Ordinal
psych: Ordinal
TSH measured: Ordinal
TSH: Ratio
T3 measured: Ordinal
T3: Ratio
TT4 measured: Ordinal
TT4: Ratio
T4U measured: Ordinal
T4U: Ratio
FTI measured: Ordinal
FTI: Ratio
TBG measured: Ordinal
TBG: Ratio
referral source: Ordinal
Condition: Nominal
Encoding Scheme:
sex: Label Encoding (Ordinal)
on thyroxine: Label Encoding (Ordinal)
query on thyroxine: Label Encoding (Ordinal)
on antithyroid medication: Label Encoding (Ordinal)
sick: Label Encoding (Ordinal)
pregnant: Label Encoding (Ordinal)
thyroid surgery: Label Encoding (Ordinal)
I131 treatment: Label Encoding (Ordinal)
query hypothyroid: Label Encoding (Ordinal)
query hyperthyroid: Label Encoding (Ordinal)
lithium: Label Encoding (Ordinal)
goitre: Label Encoding (Ordinal)
tumor: Label Encoding (Ordinal)
hypopituitary: Label Encoding (Ordinal)
psych: Label Encoding (Ordinal)
TSH measured: Label Encoding (Ordinal)
T3 measured: Label Encoding (Ordinal)
TT4 measured: Label Encoding (Ordinal)
T4U measured: Label Encoding (Ordinal)
FTI measured: Label Encoding (Ordinal)
TBG measured: Label Encoding (Ordinal)
referral source: Label Encoding (Ordinal)
Condition: One-Hot Encoding (Nominal)

Range for Numeric Variables:

Record ID: Min = 840801013, Max = 870119035, Range = 29318022

age: Min = 1, Max = 65526, Range = 65525

TSH: Min = 0.005, Max = 530.0, Range = 529.995

T3: Min = 0.05, Max = 18.0, Range = 17.95

TT4: Min = 2.0, Max = 600.0, Range = 598.0

T4U: Min = 0.17, Max = 2.33, Range = 2.16

FTI: Min = 1.4, Max = 881.0, Range = 879.6

TBG: Min = 0.1, Max = 200.0, Range = 199.9

Mean and Standard Deviation (or Variance) for Numeric Variables:

Record ID: Mean = 852947346.61, Standard Deviation = 7581968.780346589, Variance = 57486250586150.34

age: Mean = 73.56, Standard Deviation = 1183.9767180444667, Variance = 1401800.8688713466

TSH: Mean = 5.22, Standard Deviation = 24.184006144749777, Variance = 584.8661532092949

T3: Mean = 1.97, Standard Deviation = 0.8875788237425206, Variance = 0.7877961683561565

TT4: Mean = 108.70, Standard Deviation = 37.52267036706598, Variance = 1407.9507914754913

T4U: Mean = 0.98, Standard Deviation = 0.2003604411805482, Variance = 0.04014430639006391

FTI: Mean = 113.64, Standard Deviation = 41.551649606979, Variance = 1726.5395850611578

TBG: Mean = 29.87, Standard Deviation = 21.080503860189545, Variance = 444.3876429994663

Missing Values: Record ID 0

age	0
sex	307
on thyroxine	0
query on thyroxine	0
on antithyroid medication	0
sick	0
pregnant	0
thyroid surgery	0
I131 treatment	0
query hypothyroid	0
query hyperthyroid	0
lithium	0
goitre	0
tumor	0
hypopituitary	0
psych	0
TSH measured	0
TSH	842
T3 measured	0
T3	2604
TT4 measured	0

TT4	442	
T4U measured	0	
T4U	809	
FTI measured	0	
FTI	802	
TBG measured	0	
TBG	8823	
referral source	0	
Condition	0	
dtype: int64		
Outliers per column: Record ID		0
age	4	
sex	0	
on thyroxine	1240	
query on thyroxine	153	
on antithyroid medication	116	
sick	344	
pregnant	107	
thyroid surgery	134	
I131 treatment	169	
query hypothyroid	630	
query hyperthyroid	651	
lithium	93	
goitre	84	
tumor	241	
hypopituitary	2	
psych	418	
TSH measured	842	
TSH	884	
T3 measured	0	
T3	360	
TT4 measured	442	
TT4	422	
T4U measured	809	
T4U	420	
FTI measured	802	
FTI	501	
TBG measured	349	
TBG	29	
referral source	0	
Condition	2401	
dtype: int64		
Missing values after imputation: Record ID		0
age	0	
sex	0	
on thyroxine	0	
query on thyroxine	0	
on antithyroid medication	0	

sick	0
pregnant	0
thyroid surgery	0
I131 treatment	0
query hypothyroid	0
query hyperthyroid	0
lithium	0
goitre	0
tumor	0
hypopituitary	0
psych	0
TSH measured	0
TSH	0
T3 measured	0
T3	0
TT4 measured	0
TT4	0
T4U measured	0
T4U	0
FTI measured	0
FTI	0
TBG measured	0
TBG	0
referral source	0
Condition	0

dtype: int64

Normalized Data	Sample:	Record ID	age	sex	on thyroxine	query on thyroxine \
0	0.000000e+00	0.000427	0.0		0.0	
1	3.410871e-08	0.000427	0.0		0.0	
2	9.891527e-07	0.000610	0.0		0.0	
3	6.934301e-05	0.000534	0.0		0.0	
4	6.937712e-05	0.000473	0.0		0.0	

on antithyroid medication	sick	pregnant	thyroid surgery	I131 treatment \
0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0

...	TT4 measured	TT4	T4U measured	T4U	FTI measured \
0	...	0.0	0.170569	0.0	0.365741
1	...	1.0	0.210702	0.0	0.365741
2	...	0.0	0.170569	0.0	0.365741
3	...	0.0	0.170569	0.0	0.365741
4	...	0.0	0.170569	0.0	0.365741

	FTI	TBG measured	TBG	referral source	Condition
0	0.122328	0.0	0.129565	1.0	0.806452
1	0.122328	0.0	0.129565	1.0	0.806452
2	0.122328	1.0	0.054527	1.0	0.806452
3	0.122328	1.0	0.129565	1.0	0.806452
4	0.122328	1.0	0.179590	1.0	1.000000

[5 rows x 31 columns]

Jaccard Coefficient: 0.4, SMC: 0.8636363636363636

Cosine Similarity: 0.6605204689109685

